

# Skin Health Analysis App – Technical Specification

## Overview

This prototype app enables users to upload skin photos (desktop or mobile) and uses AI to screen for visual biomarkers of systemic/metabolic conditions. Advanced deep-learning models (CNNs) have shown exceptional performance in dermatology image analysis – for example, transfer-trained networks (ResNet, DenseNet, etc.) have even outperformed dermatologists on multi-class skin lesion tasks <sup>1</sup>. Deep neural networks can recognize subtle texture, color, and pattern differences in skin that are hard for humans to see <sup>2</sup>. Our app will leverage these AI capabilities: an image-processing pipeline will analyze features like hyperpigmented patches, rash patterns, dryness or yellowness, and output risk scores and tailored recommendations.

The app must be lightweight and anonymous (no user accounts). Users simply upload a JPEG/PNG image of their skin. The backend processes it with computer-vision and ML modules and returns a report page. Key high-level requirements:

- **Photo Upload:** User selects or drags an image file into a web form (HTML form or Streamlit uploader). Limit file types (JPEG/PNG) and size (e.g.  $\leq 5\text{MB}$ ).
- **AI Analysis:** Use a CNN-based model to detect specific skin signs (acanthosis nigricans, xanthelasma, rashes, dryness, etc.) and estimate their likelihoods.
- **Score Output:** For each detected condition, present a percentage-based confidence (e.g. “64% likelihood of insulin-related hyperpigmentation”). These percentages come from the model’s probabilistic outputs.
- **Recommendations:** Provide categorized advice based on each finding. For example:
  - *Lifestyle:* hydration, sleep, exercise tips.
  - *Diet:* nutritional changes (e.g. low-glycemic foods if hyperpigmentation is flagged <sup>3</sup>).
  - *Medical:* e.g. “consider discussing these signs with a dermatologist or endocrinologist,” or action items like checking blood lipid or glucose levels if warranted by findings <sup>4</sup> <sup>3</sup>.
- **Presentation:** Display the results on the frontend in a clear format (text, lists, maybe charts), ensuring readability on both desktop and mobile.
- **Anonymity:** Do not collect any personal data or require login. Images are processed in-memory or in a temporary location and then discarded; no persistent storage of user data is needed.

This document details the design and technology choices to implement the above functions on Replit.

## Functional Requirements

- **Image Upload:** A simple web interface (HTML/CSS or Streamlit) with an `<input type="file">` to upload a photo. The interface should allow mobile camera roll selection. Validations should check the file type and size.
- **Preprocessing Pipeline:** After upload, the server (Flask or Streamlit backend) reads the image. Preprocessing (using OpenCV/Pillow) resizes/crops the image and enhances features (e.g. color

conversion) for analysis. As a basic step, convert the image to a standard size (e.g. 224×224) and color space; skin-area segmentation via HSV thresholding is optional to isolate flesh tones <sup>5</sup> .

- **ML Inference:** A CNN model processes the image to score each skin biomarker. This may be multi-label (each condition separately) or multi-class. For each target feature (e.g. hyperpigmentation, dryness, rash, yellow lesions), output a probability that is converted to a percentage (e.g. 0.64 → "64%").
- **Report Generation:** Based on the model's output, generate a report. The report should list each detected feature and its percentage score. It should then offer bulleted recommendations organized by category (Lifestyle, Diet, Medical). These recommendations can be templated statements triggered by the presence of certain findings. For example, if hyperpigmentation (acanthosis nigricans) is detected, include advice to reduce sugar intake (since "higher sugar levels and increased insulin concentrations" can visibly affect skin <sup>6</sup> ). If yellowish plaques (xanthelasma) are found, note the association with lipid disorders <sup>4</sup> and suggest a cholesterol check. Always include a general medical disclaimer/advice line.
- **User Interface:** The frontend should display the report in a clean, readable layout. For example, use headers, paragraphs, and lists. An example structure might be:
- **Findings:** e.g. "Insulin-related Hyperpigmentation: 64%."
- **Recommendations:** with sub-bullets for *Lifestyle, Diet, Medical*.  
The UI should remain simple (no user account screens) and responsive (e.g. using Bootstrap or simple flex layouts) for mobile compatibility.
- **Anonymity & Security:** No authentication or data logging. Temporary images should be deleted after processing. If hosted publicly, use HTTPS to encrypt uploads. Since no PII is involved, compliance is straightforward, but a simple privacy note ("Images are not stored or shared") should be displayed.

## System Architecture

The app follows a standard web client-server model:

- **Client (Frontend):** A web page (HTML/CSS/JS) or Streamlit interface that lets the user upload an image and then displays results.
- **Server (Backend):** A Python process (Flask or Streamlit) running on Replit. It handles incoming HTTP requests, runs the ML analysis, and returns HTML or rendered output.
- **Image Analysis Module:** Once the backend receives an image, it passes the data to a computer-vision and ML pipeline (described below). This pipeline returns numerical scores.
- **Recommendation Engine:** A logic component that maps detected signs to advice text. This can be rule-based (if-then based on model outputs) using static templates in code.
- **Data Flow:**
  - User visits front page and uploads file.
  - Frontend sends HTTP POST with image to backend.
  - Backend saves image to memory or temp file.
  - Preprocessing (resize, color adjust) is applied.
  - The image is fed into the CNN model; predictions are obtained.
  - Predictions are formatted into percentages.
  - The recommendation logic uses these predictions to assemble textual advice.
  - Backend renders a response page (or Streamlit output) with the results and advice.
  - The temp image is deleted; session ends.

Replit's environment supports Python web servers out-of-the-box. For example, Replit provides a **Flask template** with hot-reload enabled <sup>7</sup>. The Replit docs note that "Flask's development server automatically reloads when you make changes to your code" <sup>7</sup>, which is convenient for this prototype. Streamlit also runs on Replit (see Replit's blog on Streamlit support), so deployment is straightforward either way.

## Image Processing Pipeline

Before inference, images undergo processing to standardize input and enhance relevant features:

1. **Read and Validate:** Load the uploaded file with Python (e.g. using Pillow `Image.open()` or OpenCV `cv2.imread()`). Verify format/size constraints.
2. **Resize:** Scale the image to a fixed resolution that the model expects (common sizes are 224×224 or 256×256 pixels). This ensures consistent input dimensions.
3. **Color Conversion (Optional):** Convert the image to a color space (e.g. HSV or YCrCb) that makes skin tones easier to isolate. The PyImageSearch tutorial suggests converting RGB→HSV and using `cv2.inRange` with predefined skin-tone HSV bounds to create a skin mask <sup>5</sup>. This mask can filter out background and focus the analysis on the skin regions.
4. **Normalization:** If using a neural network, normalize pixel values as required (e.g. scale to [0,1] or mean-zero).
5. **Augmentation (Offline):** For training a model, one would augment images (flips, rotations). For inference in the prototype, this is not needed.
6. **Data Conversion:** Convert the final image array into a tensor or array shape suitable for the model (e.g. a NumPy array shape `(1, 224, 224, 3)`).

By framing skin detection as an initial color-thresholding step, OpenCV can reduce background noise <sup>5</sup>. However, this step can be omitted if the user's photo is already focused on skin. In any case, OpenCV (or the `PIL` library) will handle loading and basic image manipulation.

## Skin Feature Classification (AI Model)

The core of the system is a convolutional neural network that identifies skin biomarkers. We can approach this as a **multi-label classification** problem: the network outputs a probability for each condition of interest. For example, labels might include "*Insulin-related Hyperpigmentation (Acanthosis)*", "*Dry/Dehydrated Skin*", "*Erythematous Rash*", "*Xanthelasma/Yellow Lesions*", etc. These labels correspond to known visual markers of conditions (e.g. acanthosis nigricans for insulin resistance, xanthelasma for lipid disorders).

**Model Selection:** Use a well-known CNN backbone with transfer learning. TensorFlow/Keras provides many pretrained models (VGG16, ResNet50, InceptionV3, MobileNet, etc.) that have been trained on ImageNet <sup>8</sup>. These models can be fine-tuned on a smaller domain-specific dataset. For prototype purposes, one might start with MobileNetV2 or EfficientNet for efficiency on limited hardware. Keras example:

```
base_model = tf.keras.applications.MobileNetV2(input_shape=(224,224,3),
                                                include_top=False,
                                                weights='imagenet')
```

These pretrained layers act as feature extractors. A few new layers (GlobalPool + Dense+Sigmoid) can be added on top for our specific labels.

**Training Data:** Ideally, one would gather a dataset of skin images labeled with these markers. In practice, for a quick prototype we might *simulate* this by training on related public datasets (e.g. ISIC lesions for pigmented vs normal skin) or use synthetic augmentation. The tech spec assumes a model can be trained or fine-tuned offline. Research has shown that deep CNNs achieve very high accuracy on skin lesion classification. For example, Rezvantab *et al.* used DenseNet, ResNet, Inception networks on a large skin dataset and found “all deep learning models outperform dermatologists (by at least 11%)” <sup>1</sup>. This underscores that using a CNN here is reasonable.

**Inference:** The final model will output sigmoid probabilities for each label. These can be multiplied by 100 for a percentage “likelihood”. For instance, if the model outputs 0.64 for the hyperpigmentation label, we say “64% chance of insulin-related hyperpigmentation.” We can set a threshold (e.g. 0.5) to decide which markers are “present,” but the report can still list all scores. Calibration or softmax is optional; for interpretability we stick with independent sigmoid scores (multiple conditions can co-occur).

No specialized object-detection is needed (we do not need bounding boxes), since each image is analyzed as a whole. If high precision is required (e.g. locating a rash region), later versions could add segmentation (U-Net) or detection (YOLO), but the prototype can treat the image globally.

## Backend Implementation

Implement the server in Python. Two main options: **Flask** or **Streamlit**.

- **Flask:** A lightweight WSGI web framework, ideal for custom endpoints <sup>9</sup>. In Flask, one would define routes such as:
  - `GET /` to render the upload form (HTML template with `<form>`).
  - `POST /analyze` to receive the uploaded image (accessed via `request.files['skin_image']`), process it, and return a rendered HTML results page. Flask templates (Jinja2) can be used to populate the report with scores and advice. This approach gives full control over HTML/CSS. Flask is well supported on Replit (Replit docs provide a Flask quickstart) <sup>7</sup>.
- **Streamlit:** An open-source framework for rapidly building data apps in pure Python <sup>10</sup>. With Streamlit, the entire app can be a single `main.py` script. A file upload widget (`st.file_uploader`) handles the image input. The same script can then run the analysis and use `st.write()`, `st.image()`, and `st.markdown()` to display text and results. Streamlit abstracts away routing and front-end code. It can display charts or images easily. It does require running `streamlit run main.py`, which on Replit may require some setup (the Replit community has a Streamlit template). Streamlit is particularly quick to iterate with and automatically updates the UI when the script changes <sup>10</sup>.

Either approach is viable on Replit. If using Flask, deploy by running `python3 main.py` (as per Replit's template) <sup>7</sup>. If using Streamlit, one can run `streamlit run main.py` (and possibly configure the run button in Replit). Both can serve an interactive web interface.

#### Libraries:

- **Web:** If Flask, install `flask`. If Streamlit, `streamlit`.
- **Image Processing:** `opencv-python` or `Pillow` (PIL).
- **ML:** `tensorflow` or `torch` (we recommend TensorFlow/Keras for easier prototyping and availability of pretrained models <sup>8</sup>).
- **Other:** `numpy`, `scikit-learn` (for any metrics).
- **Deployment:** No database needed. Static text (recommendations) can be hard-coded or stored in config files.

**Concurrency:** For a prototype, a single-threaded dev server is sufficient. If scaling, Flask can run on multiple workers, but Replit's autoscaling is beyond the prototype scope.

## Frontend Design

The front end should be minimal yet user-friendly:

- **Upload Form:** An HTML form with `enctype="multipart/form-data"`, `<input type="file">` (accepting images), and a submit button. Example snippet:

```
<form action="/analyze" method="POST" enctype="multipart/form-data">
  <input type="file" name="skin_image" accept="image/png, image/jpeg"
  required>
  <button type="submit">Analyze Skin Photo</button>
</form>
```

This can be styled simply with CSS or a framework like Bootstrap for responsiveness. The form should validate (client-side) that a file is chosen.

- **Result Display:** After submission, the page should show the report. Structure for readability: use headings and paragraphs. For example:

```
<h2>Analysis Results</h2>
<p><strong>Insulin-related Hyperpigmentation:</strong> 64% likelihood.</p>
<h3>Recommendations</h3>
<ul>
  <li><strong>Lifestyle:</strong> Stay well-hydrated and maintain a regular
  sleep schedule. Moderate exercise can help insulin sensitivity.</li>
  <li><strong>Diet:</strong> Choose whole grains and vegetables (complex
  carbs) instead of sugary foods 3. This avoids blood sugar spikes
  associated with skin issues.</li>
```

```
<li><strong>Medical:</strong> Consider discussing these findings with a
dermatologist or endocrinologist. Dark neck patches (acanthosis nigricans)
can signal pre-diabetes.</li>
</ul>
```

Each list item combines the category label and advice. Use CSS classes to style these sections with spacing and colors if desired. Ensure text is legible on mobile (e.g. using a responsive grid or scalable fonts).

If using Streamlit, the equivalent would be:

```
st.title("Analysis Results")
st.write(f"**Insulin-related Hyperpigmentation:** {score:.0%} likelihood")
st.subheader("Recommendations")
st.markdown("- **Lifestyle:** Stay hydrated, get 7-9 hours of sleep, and
exercise regularly.")
st.markdown("- **Diet:** Focus on low-glycemic foods (whole grains, vegetables)
and reduce sugar 3 .")
st.markdown("- **Medical:** Consider a check-up with a dermatologist/
endocrinologist; acanthosis nigricans may indicate insulin issues.")
```

This code will automatically render styled markdown in the app.

## Recommendations Logic

For each detected condition, predefined advice is mapped. Example mappings:

- **Hyperpigmentation (Acanthosis Nigricans):** Indicates insulin resistance risk.
  - *Lifestyle:* "Increase water intake, improve sleep quality, regular moderate exercise (e.g. 30 min/day)."
  - *Diet:* "Opt for complex carbohydrates and fiber (vegetables, whole grains) instead of high-glycemic foods <sup>3</sup> ; reduce refined sugar."
  - *Medical:* "Dark velvety skin patches often correlate with insulin resistance <sup>6</sup> . Discuss these findings with your doctor; a glucose tolerance test or A1C may be advisable."
- **Yellowish Lesions (Xanthelasma):** Suggests lipid metabolism issues.
  - *Lifestyle:* "Maintain a healthy weight and avoid tobacco."
  - *Diet:* "Limit saturated and trans fats (butter, fried foods); include foods high in omega-3 (fish, flaxseed)."
  - *Medical:* "Yellow eyelid plaques (xanthelasma) are commonly linked to high cholesterol <sup>4</sup> . Consider lipid profile testing and consulting a physician."
- **Dry/Flaky Skin:** May indicate dehydration or thyroid issues.

- *Lifestyle*: "Use moisturizers and avoid hot showers; ensure adequate hydration."
- *Diet*: "Include omega-3 fatty acids (fish, nuts) and vitamin-rich foods."
- *Medical*: "Persistent very dry skin could be a sign of hypothyroidism or other conditions; consult a dermatologist if it doesn't improve."
- **Red Rash**: Could be allergies, dermatitis, or autoimmune.
- *Lifestyle*: "Identify and avoid irritants (new soaps, fabrics); use gentle skincare."
- *Diet*: "Anti-inflammatory diet: include fruits, vegetables, and omega-3s; reduce dairy/spices if they trigger you."
- *Medical*: "If rash persists or worsens, consider seeing a dermatologist. Certain patterns (e.g. butterfly rash) may warrant rheumatology evaluation."

Each bullet should be concise (3–5 sentences total per category recommended in guidelines) and friendly but not alarmist. Citation [18] is used above to support diet advice on glycemic index. Citation [29] supports linking yellow lesions to lipid issues. These references assure our recommendations are grounded in medical observations.

## Technology Stack

- **Programming Language**: Python 3.x.
- **Image Processing**: [OpenCV](#) (`cv2` in Python) for tasks like color-space conversion and resizing <sup>5</sup>. Optionally Pillow (`PIL`) for simple image I/O.
- **Machine Learning**: TensorFlow 2.x with Keras API. Keras includes many pretrained CNNs suitable for transfer learning <sup>8</sup>. Example choices: MobileNetV2 (lightweight), EfficientNet (accurate but heavier), or ResNet50. For prototyping, MobileNetV2 is fast on limited CPU/GPU. PyTorch is an alternative, but TensorFlow's Keras has easy model loading (e.g. `tf.keras.applications.ResNet50`).
- **Backend Framework**: *Flask* (preferred) or *Streamlit*. Flask is a microframework for building web applications <sup>9</sup>; it handles routes, form data, and template rendering. Streamlit is a high-level app framework for data apps <sup>10</sup>; it can eliminate front-end coding but offers less flexibility. Both run on Replit.
- **Frontend**: Standard web technologies (HTML5, CSS3, JavaScript). Use semantic HTML for form and results. For styling and responsiveness, a CSS framework like Bootstrap or Bulma is recommended. No front-end framework (React/Vue) is needed for this simple UI. Streamlit, if used, handles UI elements via Python.
- **Deployment**: Host on Replit. Replit's built-in server (auto-reloaded Flask dev server) simplifies development <sup>7</sup>. There's no separate database; if needed, a simple JSON file or Python dict can hold static recommendation texts.
- **Version Control**: Use Replit's Git integration or link to GitHub for code history.

## Implementation Notes

- **Flask Example**: Use `flask` and `flask_cors` (optional) to create endpoints. A basic route might be:

```

from flask import Flask, request, render_template
app = Flask(__name__)

@app.route('/', methods=['GET'])
def index():
    return render_template('upload.html')

@app.route('/analyze', methods=['POST'])
def analyze():
    file = request.files['skin_image']
    img = preprocess(file)          # custom function
    scores = model.predict(img)      # yields probabilities
    report = build_report(scores)    # custom logic
    return render_template('result.html', report=report)

```

Templates `upload.html` and `result.html` use Jinja2 to insert dynamic content. After returning the result, delete or close the uploaded file.

- **Streamlit Example:** A single script `main.py`:

```

import streamlit as st
from PIL import Image

st.title("Skin Health Analysis")
uploaded = st.file_uploader("Upload a skin photo",
type=["jpg", "jpeg", "png"])
if uploaded:
    img = Image.open(uploaded)
    st.image(img, caption="Uploaded Image", use_column_width=True)
    scores = model.predict(preprocess(img))
    st.write(f"**Hyperpigmentation (Insulin-related):**")
    {scores['hyperpig']*100:.1f}%")
    # ... other scores
    st.subheader("Recommendations")
    st.write("- **Lifestyle:** ...")
    # ...

```

Streamlit auto-refreshes when code changes, and widgets simplify the workflow.

- **Model Serving:** Load the trained model once at startup (global scope) to avoid reloading on each request. This keeps inference fast. If using TensorFlow, call `model = tf.keras.models.load_model(...)` outside of route handlers.
- **Error Handling:** Check for corrupted uploads. If preprocessing or prediction fails, return a user-friendly error. Add a basic try/except around analysis.



- **Testing:** Test with sample images representing each condition. Ensure percentages are plausible (e.g. simulating 0–100%).

## Privacy and Security

- **No Login:** The app is anonymous. Do not implement user authentication or sessions beyond the current upload.
- **Transient Data:** Handle images in RAM or temp storage; delete immediately after use. Do not write images or results to a persistent database.
- **Communication:** If deployed beyond Replit's preview, use HTTPS. (Replit subdomains use TLS automatically.)
- **Data Use:** Display a brief disclaimer like “This tool provides probabilistic analysis and is not a medical diagnosis. Consult a healthcare professional for concerns.”

## References

- AI in Dermatology: Deep convolutional networks greatly improve image-based diagnosis efficiency, recognizing subtle lesion features <sup>2</sup>. In one study, several pretrained CNNs (DenseNet, ResNet, etc.) achieved accuracy surpassing human dermatologists by 11% <sup>1</sup>.
- Pretrained Models: Keras includes 19+ state-of-the-art models (VGG, ResNet, Inception, etc.) trained on ImageNet <sup>8</sup>, which can be fine-tuned for skin analysis.
- Skin Detection (OpenCV): A simple method is converting the image to HSV and applying `cv2.inRange` with skin-tone thresholds to create a binary mask <sup>5</sup>.
- Web Frameworks: Flask is “a lightweight WSGI web application framework” suited for Python web apps <sup>9</sup>. Streamlit allows quick creation of data apps in Python <sup>10</sup>. Replit provides templates and auto-reload for Flask <sup>7</sup>.
- Medical Indicators: High sugar/insulin levels manifest in skin (e.g. dark velvety patches) <sup>6</sup>; diet advice recommends low-glycemic foods <sup>3</sup>. Yellow eyelid xanthelasma indicates lipid disorders <sup>4</sup>. These inform our diet/medical recommendations.

---

<sup>1</sup> <sup>2</sup> Artificial Intelligence in Dermatology Image Analysis: Current Developments and Future Trends - PMC

<https://pmc.ncbi.nlm.nih.gov/articles/PMC9693628/>

<sup>3</sup> <sup>6</sup> How Sugar is Ruining Your Skin - Spring Street Dermatology

<https://springstderm.com/how-sugar-is-ruining-your-skin/>

<sup>4</sup> Xanthomas

<https://dermnetnz.org/topics/xanthoma>

<sup>5</sup> Tutorial: Skin Detection Example using Python and OpenCV

<https://pyimagesearch.com/2014/08/18/skin-detection-step-step-example-using-python-opencv/>

<sup>7</sup> Replit Docs

<https://docs.replit.com/getting-started/quickstarts/flask-app>

<sup>8</sup> Image Classification using Pre-Trained ImageNet Models in TensorFlow & Keras

<https://learnopencv.com/image-classification-pretrained-imagenet-models-tensorflow-keras/>

9 Welcome to Flask — Flask Documentation (3.1.x)

<https://flask.palletsprojects.com/en/stable/>

10 Streamlit • A faster way to build and share data apps

<https://streamlit.io/>